

AI local models, last 30 days: the year running AI on your own machine stopped being a hobby


An open model overtook the closed frontier, Apple shipped a first-party local-agent stack, and the best-watched videos all answered the same question — how do I run capable AI without the cloud or the bill?

Topic: "AI local models" · 5 top videos by views · window 2026-03-24 → 2026-06-22 · engagement-ranked, Shorts & live excluded · generated by last30daysYT

TL;DR

- **An open model just took the lead.** ZAI's GLM 5.2 is "by far the best open model," reportedly beating the best GPT and Gemini on multiple benchmarks. [\[00:00\]](#)
- **Apple made local agents first-class.** A WWDC26 session runs a full agentic loop on-device via MLX — "no cloud, no API keys." [\[00:07\]](#)
- **Local agents are "a new category."** The framing that pulled 322k views: agent = LLM brain + text-file memory + tools + a cron "heartbeat." [\[00:00\]](#)
- **The motivation is money.** The most-watched video (552k) opens flatly: "AI pricing is outrageous." [\[00:00\]](#)
- **It runs on a Raspberry Pi now.** Gemma 4's smallest model answers real prompts on an 8GB Pi 5 — slowly, but usably. [\[08:48\]](#)

ENGAGEMENT — VIEWS PER FEATURED VIDEO



THEME 1 — THE MODELS GOT GOOD ENOUGH

An open model overtook the closed frontier

The single biggest story is ZAI's **GLM 5.2**. The reviewer's claim is blunt: "We have a new number one open-source model and the gap is insane... it even beats the best GPT and the best Gemini across multiple benchmarks." [\[00:00\]](#) The demos are the evidence — in single prompts (with light nudging) it built an interactive 3D digital twin of Earth, an exploded-view V8 engine, and a mechanical watch with turning gears.

"The thing I like about GLM is it requires very minimal handholding. It just works right out of the box most of the time." [\[06:47\]](#)

Why it matters: the reviewer benchmarks it against the closed frontier (Claude Fable 5) and it lands in the same conversation — for a weight you can download. The headline closing the cloud-vs-local gap is a model, not a tool.

...and it now fits on a \$80 computer

At the other end of the scale, a developer runs **Gemma 4**'s smallest model (E2B, ~4.5 GB, ~4B params) on a Raspberry Pi 5 with 8 GB of RAM. [\[01:03\]](#) It's Apache-2.0 licensed with a 128k context and native audio/image support. [\[02:37\]](#) A sorting task took ~6 minutes

including reasoning — slow, but it produced two implementations and explained the trade-offs.

"It's absolutely possible, and the results are usable, at least for certain scripts or automation tasks." [\[08:48\]](#)

Why it matters: "local" no longer requires a GPU rig — the floor is now a board you can hide behind a monitor.

THEME 2 — THE AGENT LAYER MATURED

Apple shipped a first-party local-agent stack

The most consequential release this month is institutional: a WWDC26 session walks through running the full **agentic loop** — user → agent → model → tools, cycling until the task is done — entirely on a Mac via MLX. [\[01:10\]](#) It's a clean four-layer stack: MLX → MLX-LM → an OpenAI-compatible MLX-LM Server → any agent (Xcode, OpenCode, Pi). [\[02:46\]](#)

"No cloud, no API keys, just your hardware doing the work." [\[00:07\]](#)

The hardware angle is concrete: the M5's Neural Accelerators make matrix multiply ~4× faster than M4, which maps almost directly to prompt-processing speed [\[06:25\]](#) — the bottleneck for agents, which churn hundreds of thousands of tokens of tool output. For models too big for one machine (DeepSeek's 1.6T params need >800 GB), MLX shards across Macs over Thunderbolt. [\[07:58\]](#) Why it matters: when the platform vendor ships the stack, local agents move from enthusiast project to supported workflow.

"Local AI agents are a new category of AI product"

The most useful mental model came from a 26-minute primer. Its anatomy of an agent is refreshingly demystifying: pick where it lives, give it a communication channel (Telegram/Discord), a brain (the LLM), memory, tools, and a "heartbeat" (a cron schedule so it acts unprompted). [\[02:35\]](#)

"Memory is literally just a bunch of text files... you just dump in there everything that you want it to know." [\[06:44\]](#)

The honest part is the safety warning: people who gave an agent access to everything on their main laptop had it delete emails — so she runs hers on a wiped, dedicated old MacBook.

[\[03:36\]](#) Why it matters: the gap between "I have a local model" and "I have a useful local agent" is design, and this is the clearest map of it.

THEME 3 — THE REALITY CHECK: HARDWARE & COST

The most-watched video is a VRAM lesson, and that's the tell

552k people watched a 45-minute local-coding setup whose real subject is hardware literacy. The opener names the motive — "AI pricing is outrageous" [\[00:00\]](#) — then teaches the three numbers that decide everything: parameters, context, and VRAM. The key constraint: a model loads into your GPU's VRAM, and whatever doesn't fit overflows to system RAM and slows down. [\[04:07\]](#)

"The key takeaway: how much RAM your graphics card has, as well as how much RAM your computer overall has." [\[04:38\]](#)

It also demystifies **quantization** (Q4/Q6/Q8) — rounding the weights to shrink the model, with Q4 a good starting point and roughly half the size of Q8. [\[07:45\]](#) Why it matters: the demand signal is loud — the single biggest audience this month went to "how do I actually fit this on my machine and stop paying per token."

WORTH WATCHING

- [New #1 open-source AI model is here!](#) — AI Search · 407k views. The GLM 5.2 demos; start here for "how good is open now."
- [WWDC26: Run local agentic AI on the Mac using MLX](#) — Apple Developer · 389k views. The most authoritative how-it-works, 13 min.
- [The Best Local Agentic Coding Workflow](#) — Web Dev Simplified · 552k views. The deepest hands-on; the VRAM/quantization masterclass.

ON THE RADAR

Three threads to track next month: **independent benchmarks for GLM 5.2's "beats GPT/Gemini" claim** (so far it's the reviewer's own tests); **whether M5 / MLX makes on-device agents fast enough to replace cloud coding day-to-day**; and **the agent-safety story** — the "it deleted my emails" warning suggests sandboxing, not capability, is the next bottleneck.

Built by **last30daysYT** from YouTube transcripts (via yt2md). Topic "AI local models", top 5 by view count, last 30 days (extended to a 90-day window to fill the set). Every claim links to the exact moment in the source video. Engagement is a prefilter — substance was judged from the transcripts; all 5 cleared the bar.